Open in app ↗

Why Andrej Karpathy's "LLM Wiki" is the Future of Personal Knowledge

How to stop rediscovering information from scratch and let AI automatically compile, maintain, and compound your second brain.

5 min read · Apr 6, 2026

evoailabs

Follow

Listen Share ••• More

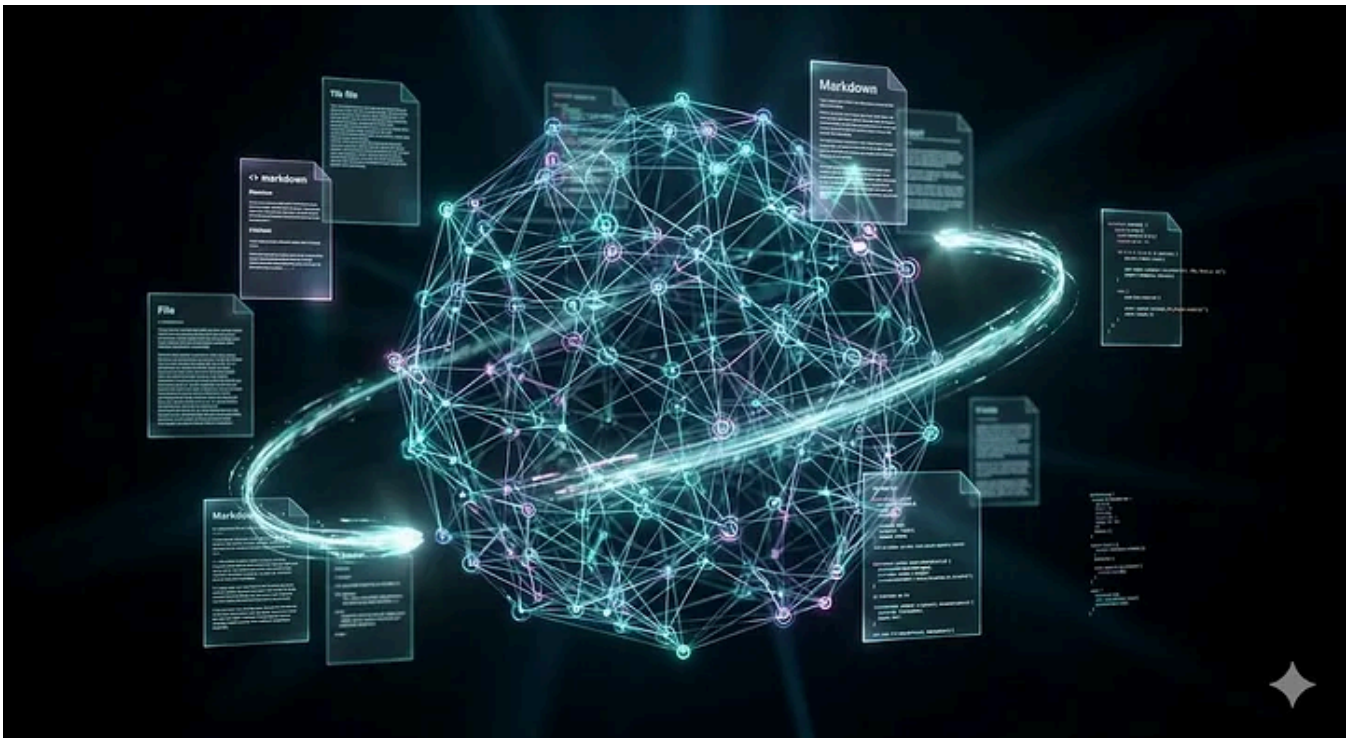

In April 2026, Andrej Karpathy (OpenAI co-founder and former Tesla AI Director) caused a stir in the AI community when he dropped a GitHub Gist titled simply "LLM Wiki". It wasn't a new app, a library, or a SaaS product. It was an "idea file" — a conceptual pattern designed to be copy-pasted into an LLM agent like Claude Code or OpenAI Codex.

https://x.com/karpathy/status/2040470801506541998

Here is a deep dive into the core ideas behind the LLM Wiki, the problems it solves, how it works, and how the broader ecosystem is already building upon it.

The Oldest Problem: Knowledge Maintenance

Currently, the default way we interact with our documents is through **Retrieval-Augmented Generation (RAG)**. You upload a bunch of PDFs, ask a question, and the LLM searches for relevant text chunks to synthesize an answer. But RAG has a critical flaw: **there is no accumulation**. If you ask a complex question that requires synthesizing five different documents, the LLM has to find and piece together those fragments from scratch. If you ask a similar question tomorrow, it does the exact same work again. Nothing is built up.

On the flip side, humans have tried manually building linked knowledge bases for decades (e.g., the Zettelkasten method, Notion wikis, Obsidian vaults). But human-maintained wikis inevitably fail because the **bookkeeping burden** — updating cross-references, tagging, and noting contradictions — grows much faster than the value.

What is the "LLM Wiki"?

The LLM Wiki is a design pattern that inserts an LLM-maintained, compounding layer of markdown files between you and your raw source materials.

Instead of searching raw documents on the fly, the **LLM pre-compiles them**. It reads your sources, extracts the concepts, and permanently weaves them into an interlinked markdown wiki.

As Karpathy perfectly frames it: **"Obsidian is the IDE, the LLM is the programmer, the wiki is the codebase."** You rarely write the wiki yourself. You explore and ask questions, while the LLM does the tedious grunt work of maintaining the database.

How It Works: Architecture & Operations

The LLM Wiki framework is built on a simple three-layer architecture:

1. **Raw Sources (Immutable):** Your curated PDFs, articles, meeting transcripts, and images. The LLM reads these but never modifies them.
2. **The Wiki (LLM-Maintained):** A directory of markdown files (summaries, concept pages, timelines) completely maintained by the LLM.
3. **The Schema (The Rules):** A configuration file (e.g., CLAUDE.md) that tells your specific agent how to structure the wiki, how to ingest new files, and how to format answers.

To keep this ecosystem alive, the agent executes three main operations:

- **Ingest:** You drop a new article into the raw folder. The agent reads it, writes a summary, updates 10–15 related concept pages with new insights, adds backlinks, and logs the action.
- **Query:** You ask a question. The agent reads the wiki's central index.md, navigates to the right pages, and gives you an answer. If you discover a new connection during your chat, the LLM files that insight back into the wiki as a new page.
- **Lint:** Periodically, you ask the LLM to "health-check" the wiki. It hunts for broken links, stale claims, contradictions, and orphan pages.

How It Can Be Used

Because it is just a set of instructions applied to standard Markdown files, the use cases are virtually limitless:

- **Personal Growth:** Feed it journal entries, health data, and podcast notes to build a structured, evolving picture of your psychology and goals.
- **Academic/Deep Research:** As you read papers over months, the LLM incrementally builds out an evolving thesis, linking methodologies and noting where researchers contradict each other.

- **Book & Hobby Tracking:** Drop in book chapters as you read, and watch the LLM build a Tolkien-style fan wiki mapping characters, locations, and plot lines.
- **Business & Teams:** Feed it Slack threads, customer calls, and PRs. New team members can instantly browse an up-to-date internal wiki that nobody had to manually write.

The Ecosystem: Similar Projects Solving the Same Problem

Karpathy's gist didn't happen in a vacuum. It gave a name to a massive shift in how developers are treating AI and knowledge. Here is an analysis of similar projects and spin-offs solving the RAG bottleneck:

1. Waykee Cortex (Hierarchical Team Knowledge)
While Karpathy's wiki relies on flat indexing, *Waykee* focuses on teams. It utilizes a strict hierarchical inheritance model (a specific UI screen inherits context from its module, which inherits from the system). It also uniquely combines the "Knowledge" layer (what exists) with the "Work" layer (tasks, bugs, milestones), allowing issues to inherit dual-context automatically.

2. Sage-Wiki (The Pipeline Approach)
Built by developer *xoai*, this implementation treats the LLM less like a conversational agent and more like a strict compiler (similar to make). It uses a 5-step incremental pass (diff -> summarize -> extract -> write -> images) and enforces a strict typed-entity system (e.g., specifying if a node "is-a" or "contradicts" another) to prevent the LLM from creating duplicate concepts.

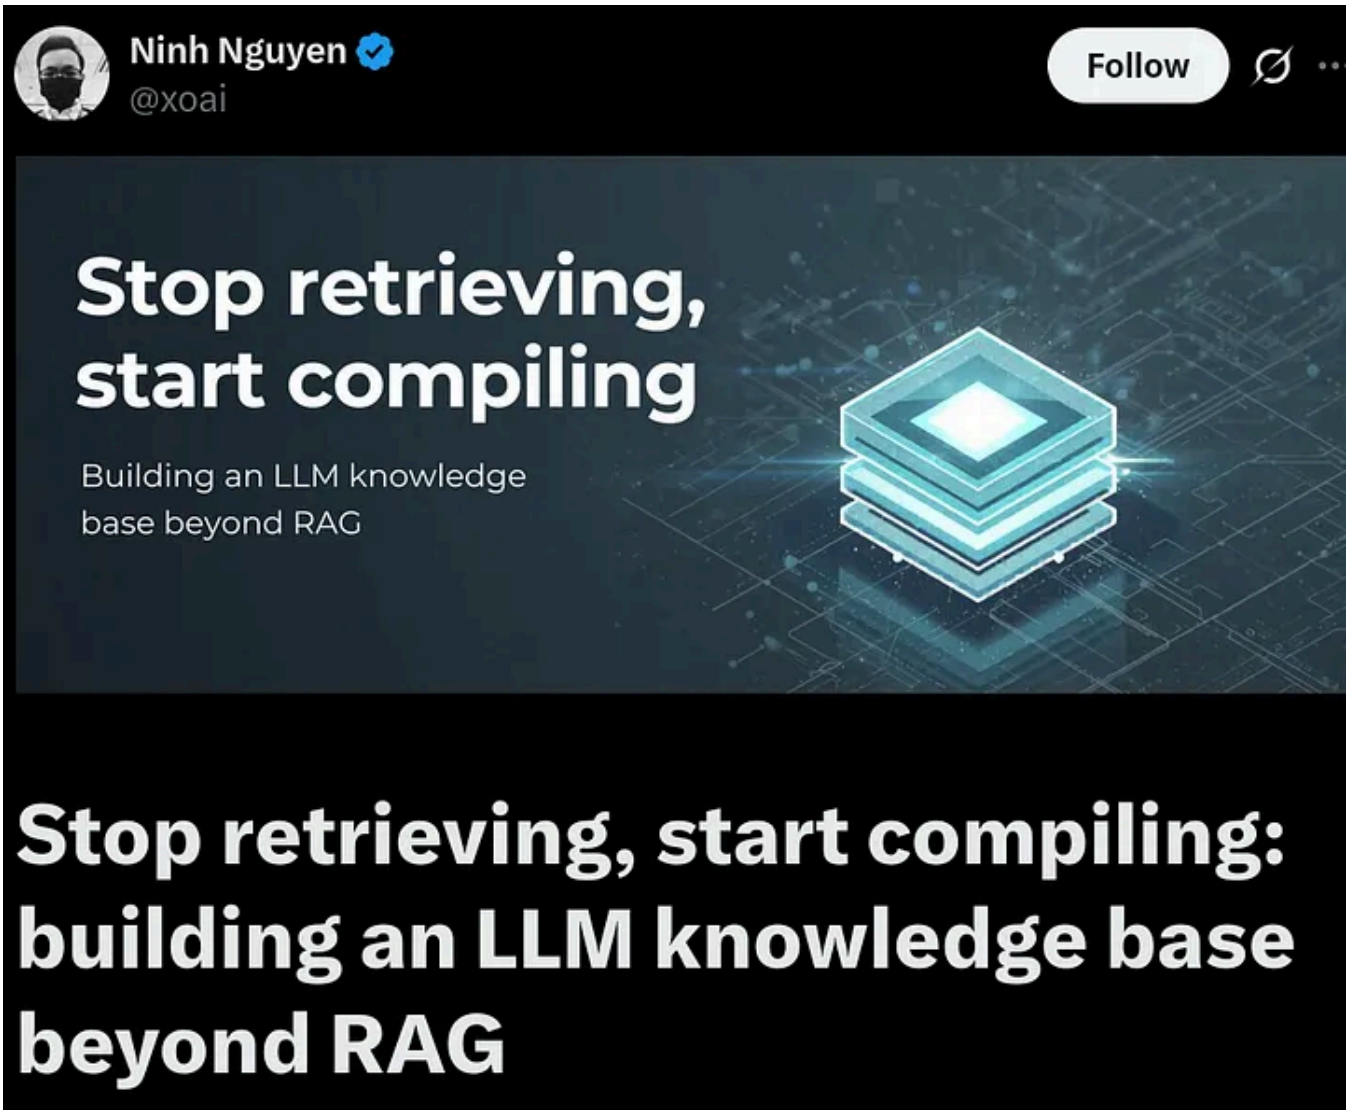

https://x.com/xoai/status/2040936964799795503

3. Thinking-MCP (Capturing Mental Models)
Rather than tracking factual data, *Thinking-MCP* captures *how you think*. It scans conversation transcripts to map your heuristics, tensions, and decision-making rules. Instead of a permanent wiki, its graph relies on "node decay" — meaning core values persist, but fleeting ideas fade over time, mirroring the live state of a human brain.

4. ELF (Eli's Lab Framework)
Designed specifically for the rigorous demands of scientific research, ELF mixes the PARA organization method with wiki architecture. It uses a "base-delta protocol" for incremental experiments, ensuring total data traceability while minimizing researcher documentation fatigue.

5. qmd (Local Markdown Search)
As wikis scale beyond a few hundred files, a simple index.md file isn't enough. Built by Shopify CEO Tobi Lütke, *qmd* solves this by acting as a local search engine over

markdown files, using a hybrid of BM25 and vector search. With an MCP server, the LLM can shell out to *qmd* natively to fetch information across massive personal wikis.

The LLM Wiki pattern marks a fundamental maturity in how we use AI. We are moving away from treating LLMs purely as search engines or text generators, and finally starting to use them **as tireless librarians and system maintainers.** By shifting the workload of documentation to the AI, we finally unlock the dream of a compounding "Second Brain" that actually takes care of itself.

The viral story continues...

The Most Important Employee at a Japanese Firm Was a Markdown File

How a Japanese tax accountant uses CLAUDE.md to serve 60 clients and what it reveals about future of knowledge work.

Kaitai Dong Follow 11 min read · 23 hours ago

https://ai.gopubby.com/the-most-important-employee-at-a-japanese-firm-was-a-markdown-file-8f079d9c4d19

Second Brain | Andrej Karpathy | LLM | Wiki

Follow

Written by evoailabs

934 followers · 37 following

Tech/biz consulting, analytics, research for founders, startups, corps and govs.

Responses (6)

Christian Broberg

What are your thoughts?

Darrell
Apr 12

The idea of an LLM-backed personal wiki sounds great in theory, but I keep wondering how it holds up once your notes get messy and inconsistent over time. Have you tried it long enough to see if it actually stays useful or starts drifting?

6 1 reply Reply

Mehmet Zaim
Apr 11

Not a plain Wiki (such as MediaWiki) that has a taxonomical structure but a Wiki that has an ontological structure (like Semantic MediaWiki) may be more helpful. In this case knowledge evolves from static text into a dynamic Knowledge Graph. Here... more

6 Reply

Matt Franz
Apr 13

Interesting, I wrote something like that here, initially as an experiment to learn Golang TUI
https://github.com/mdfranz/code-gehirn

2 Reply

See all responses

More from evoailabs

evoailabs

Self-Evolving Agents: Open-Source Projects Redefining AI in 2026

From Static Prompts to Recursive Mastery: How the Open-Source Community is Building AI that Teaches Itself

Mar 25 35 1

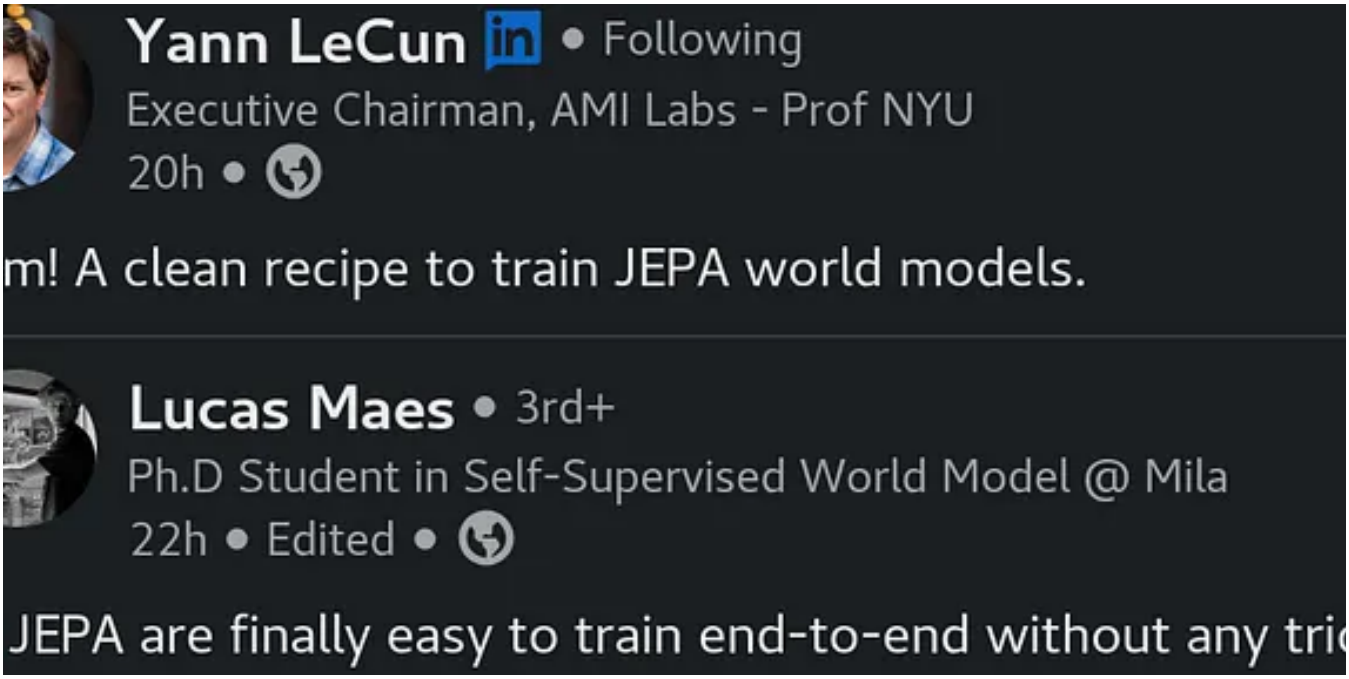

evoailabs

LeWorldModel: 15 Million Parameters to Map the World

A deep dive into the new JEPA architecture that brings common-sense physics to AI — trained on a single GPU

Mar 24 20

evoailabs

OpenClaw, NanoBot, PicoClaw, IronClaw, ZeroClaw, NullClaw: This *Claw Craziness Is Continuing...

Agent space is booming, but raw metrics don't matter: an agent's intelligence and security outweigh lines of code, language choice, and...

Feb 15 181 1

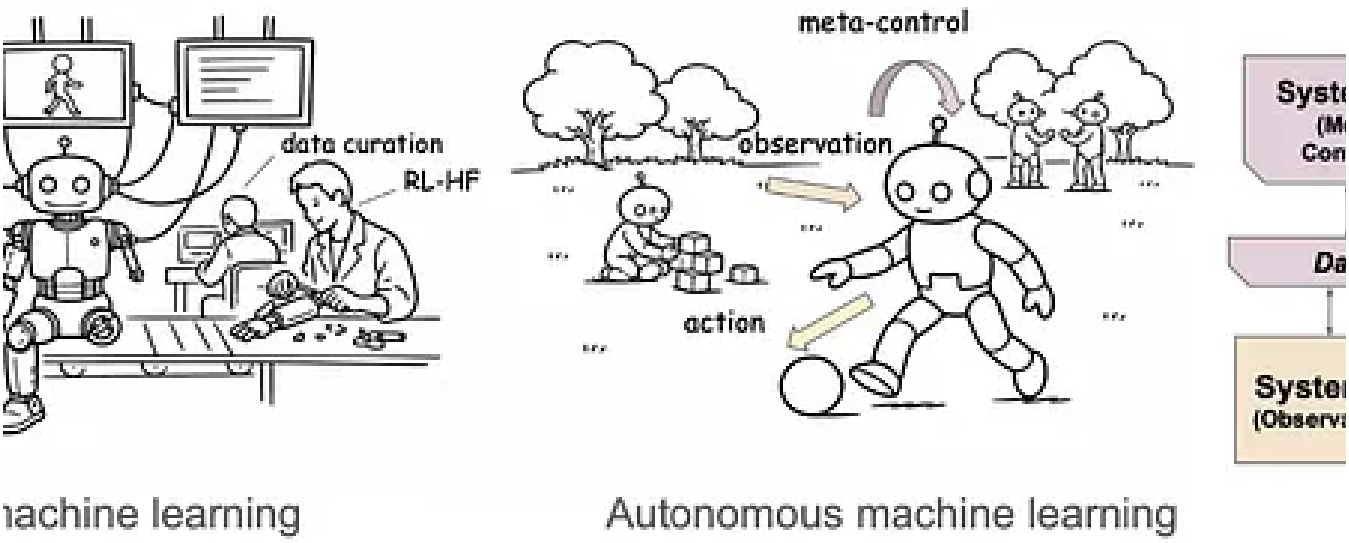

evoailabs

The “Street-Smart” Revolution: LeCun and Malik’s Manifesto for Autonomous AI

Beyond the LLM Dead-End: Why “System M” is the Key to True Machine Intelligence

Mar 18 36 2

See all from evoailabs

Recommended from Medium

In Vibe Coding by Alex Dunlop

I Cut Claude Code's Output Tokens by 75%. Why Did Nobody Tell Me?

Fewer tokens. Somehow more accurate. There's a paper.

Apr 11 1.2K 16

E2B	2.3B	4GB	85	6.1/10	6.5/10	On-device, edge
E4B	4.5B	8GB	38	7.4/10	7.8/10	Local dev, laptops
26B MoE	**3.8B active**	**18GB**	**42**	**9.1/10**	**8.9/10**	**Most use cases**

In Towards AI by Chew Loong Nian - AI ENGINEER

I Tested All 4 Gemma 4 Models: The 26B One Is Cheating (In the Best Way)

Google dropped Gemma 4 this week under Apache 2.0 — free to use, free to commercialize, no strings attached. Four model sizes landed at...

Apr 4 240 4

Balu Kosuri

I Turned Andrej Karpathy's Autoresearch Into a Universal Skill

By Balasubramanyam Kosuri

Mar 21 506 9

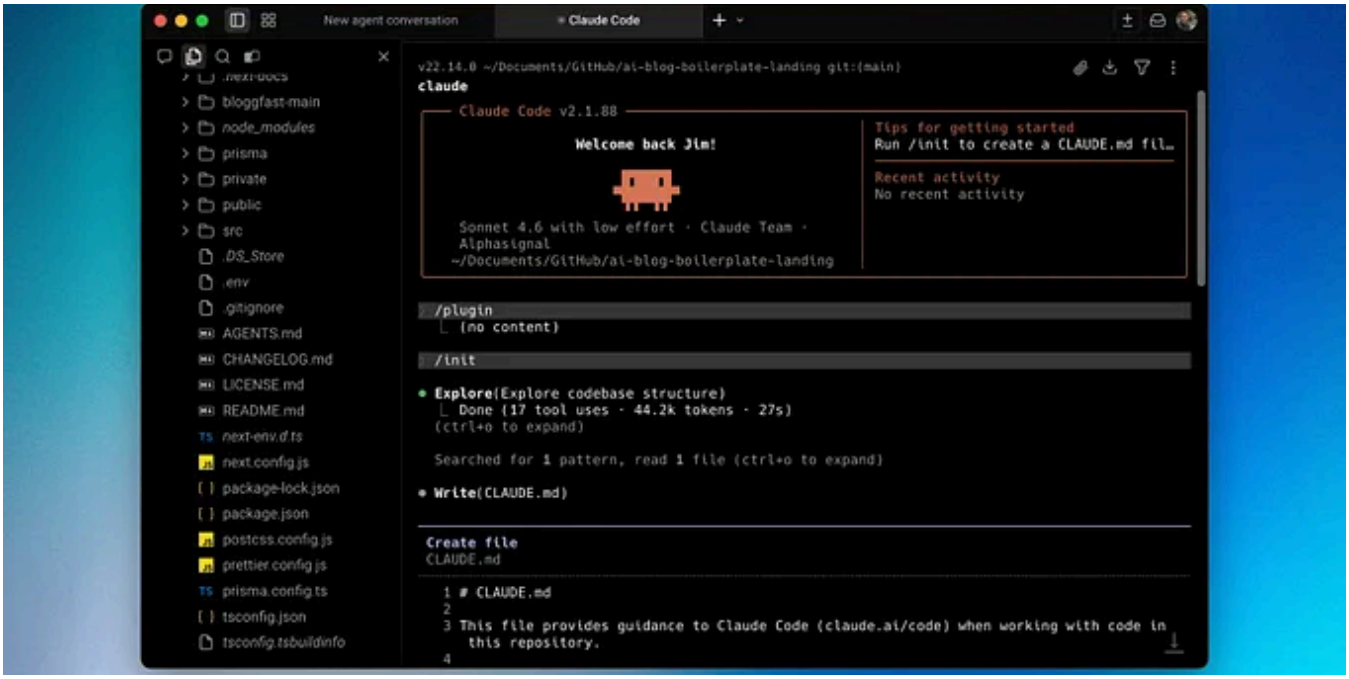

In Generative AI by Jim Clyde Monge

My Ultimate Claude Code Setup

Here's how I do my Claude Code setup that 10x my agentic development productivity.

6d ago 291

In AI Advances by Yanli Liu

I Ran Codex and Claude Side by Side. Here's What I Found.

The setup takes 5 minutes. The compliance question takes longer.

Apr 12 506 4

In Data Science Collective by Gao Dalie ( )

Gemma 4 + Turboquant + RAG: Better OCR & Self-Hosted

If you don’t have a Medium subscription, use this link to read the full article: Link

Apr 8 406 1

See more recommendations